

# UNIX Programming (18CS56)

## Module 5

### (I) Signals

- Signals are triggered by events and posted on a process to notify it that something has happened and requires some action.
- The event can be generated from a process, a user or the UNIX kernel.
- Parent and child processes can send signals to each other for process synchronization.
- Signals are the software version of hardware interrupts.
- Signals are defined as integer flags.
- `<signal.h>` header depicts the list of signals defined for a UNIX system.
- The following table lists some of the important signals used in UNIX systems –

Signal Name	Use	Generates core file at default?
<b>SIGALRM</b>	Alarm timer time-outs. Generated by <i>alarm()</i> API.	No
<b>SIGABRT</b>	Abort process execution. Generated by <i>abort()</i> API.	Yes
<b>SIGFPE</b>	Illegal mathematical operation	Yes
<b>SIGHUP</b>	Controlling terminal hang-up	No
<b>SIGILL</b>	Execution of an illegal machine instruction	Yes
<b>SIGINT</b>	Process interruption. Generated by the Delete or Ctrl+C keys.	No
<b>SIGKILL</b>	Kill a process. Generated by <i>kill -9 &lt;pid&gt;</i> command.	Yes
<b>SIGPIPE</b>	Illegal write to a pipe	Yes
<b>SIGQUIT</b>	Process quit. Generated by Ctrl+\ keys.	Yes
<b>SIGSEGV</b>	Segmentation fault. Generated by de-referencing a NULL pointer.	Yes
<b>SIGTERM</b>	Process termination. Generated by <i>kill &lt;pid&gt;</i> command.	Yes
<b>SIGUSR1</b>	Reserved to be defined by users	No
<b>SIGUSR2</b>	Reserved to be defined by users	No
<b>SIGCHLD</b>	Sent to a parent process when its child process has terminated	No
<b>SIGCONT</b>	Resume execution of a stopped process	No
<b>SIGSTOP</b>	Stop a process execution	No

<b>SIGTTIN</b>	Stop a background process when it reads from its controlling terminal	No
<b>SIGTTOU</b>	Stop a background process when it writes to its controlling terminal	No
<b>SIGTSTP</b>	Stop a process execution by the Ctrl+Z keys	No

- When a signal is sent to a process, it is *pending* on the process to handle it.
- Process can react to pending signals in one of three ways:
  - Accepts the default action of the signal. Usually terminates the process.
  - Ignore the signal. Signal will be discarded, and it has no effect on recipient process.
  - Invoke a user-defined function. The function is known as signal handler routine and signal is said to be *caught* when this function is called.

If the function finishes its execution without terminating the process, the process will continue execution from the point it was interrupted by the signal.

- Process may set up per-signal handling mechanisms – ignores some signals, catches some other signals and accepts default action from the remaining signals.
- Process may change handling of certain signals in its course of execution.
- Signal is said to have been *delivered* if it has been reacted to by the recipient process.
- The default action for most signals is to terminate a recipient process.
- Some signals will generate a core file for the aborted process – users can trace back the state of the process when it was aborted.

These signals are usually generated when there is an implied program error in the aborted process.

- Most signals can be ignored or caught except SIGKILL and SIGSTOP signals.
- Companion signal to SIGSTOP is SIGCONT, which resumes a process execution after it has been stopped.
- A process is allowed to ignore certain signals so that it is not interrupted while doing certain mission-critical work.
- Signal handler function cleans up the work environment of a process before terminating the process gracefully.

## (II) Kernel Support for Signals

- UNIX System V.3 – each entry in kernel process table slot has array of signal flags, one for each signal defined in the system.
- When a signal is generated for a process, kernel will set the corresponding signal flag in the process table slot of the recipient process.
- If recipient process is asleep, kernel will awaken the process by scheduling it.

- When recipient process runs, kernel will check the process U-area that contains an array of signal handling specifications.
- Each entry of the array corresponds to a signal defined in the system.
- Kernel will consult the array to find out how the process will react to the pending signal.
- UNIX System V.2 – when a signal is caught, kernel will first reset the signal handler in the recipient process U-area, then call the user signal handling function specified for that signal.
- If multiple instances of a signal are being sent to a process at different points, the process will catch only the first instance of the signal.
- All subsequent instances of the signal will be handled in the default manner.
- If a process is continuously catching multiple occurrences of a signal, it must reinstall the signal handler function every time the signal is caught.
- In the time between signal handler invoking and re-establishment of signal handler method, another instance of signal may be delivered to the process. This leads to race condition.
- To solve the unreliability of signal handling in UNIX System V.2, BSD UNIX 4.2 and POSIX.1 use an alternate method.
- When a signal is caught, the kernel does not reset the signal handler, so there is no need for the process to re-establish the signal handling method.
- The kernel will block further delivery of the same signal to the process until the signal handler function has completed execution.
- This ensures that signal handler function will not be invoked recursively for multiple instances of the same signal.

### (iii) signal API

- *signal* API is used to define the per-signal handling method.
- Prototype of the *signal* API is:

```
#include<signal.h>

void (*signal(int signal_num, void (*handler)(int)))(int);
```

- *signal\_num* – signal identifier such as SIGINT or SIGTERM, as defined in *<signal.h>*
- *handler* – function pointer of a user-defined signal handler function.

- The following example attempts to catch the SIGTERM signal, ignores the SIGINT signal, and accepts the default action of the SIGSEGV signal. The *pause* API suspends the calling process until it is interrupted by a signal and the corresponding signal handler does a return.

```
#include<iostream.h>
#include<signal.h>
//signal handler function
void catch_sig(int sig_num){
    signal(sig_num, catch_sig);
    cout << "catch_sig: " << sig_nm;
}
//main
int main(){
    signal(SIGTERM, catch_sig);
    signal(SIGINT, SIG_IGN);
    signal(SIGSEGV, SIG_DFL);
    pause(); //waits for signal interruption
}
```

- SIG\_IGN and SIG\_DFL are manifest constants defined in <signal.h>

```
#define SIG_DFL void (*)(int)0
#define SIG_IGN void (*)(int)1
```

- SIG\_IGN – specifies that a signal must be ignored. If the signal is generated to the process, it will be discarded without any interruption of process.
- SIG\_DFL – specifies that the default action of a signal must be accepted.
- Return value of *signal* API – previous signal handler for a signal. It can be used to restore the signal handler for a signal after it has been altered.

```
#include<signal.h>

int main(){
    void (*old_handler)(int) = signal(SIGINT, SIG_IGN);
    //do mission critical processing

    //restore previous signal handling
    signal(SIGINT, old_handler);
}
```

- UNIX System V.3 and V.4 support *sigset* API.

```
#include<signal.h>

void (*sigset(int signal_num, void (*handler)(int)))(int);
```

- Arguments and return value of *sigset* are same as that of *signal*.
- Both functions set signal handling methods for any named signal.
- *signal* API is unreliable – race condition due to handling multiple instances of signal. *sigset* API is reliable – one of the multiple instances is handled, whereas remaining instances are blocked.

#### (IV) Signal Mask

- A signal mask defines which signals are blocked when generated to a process.
- The blocked signal depends on recipient process to unblock it and handle accordingly.
- If a signal has to be specified to be ignored or blocked, it is implementation-dependent on whether such a signal will be discarded or left pending when it is sent to the process.
- A process initially inherits the parent's signal mask when it is created, but any pending signals for the parent process are not passed on.
- A process may query or set its signal mask using *sigprocmask* API:

```
#include<signal.h>

int sigprocmask(int cmd, const sigset_t *new_mask, sigset_t *old_mask);
```

- *new\_mask* – defines a set of signals to be set or reset in a calling process signal mask.
- *cmd* – specifies how the *new\_mask* value is to be used by the API.
- Returns: 0 if it succeeds, -1 if it fails.
- Possible failure - *new\_mask/old\_mask* are invalid addresses.
- The following table lists the various values that can be assigned to *cmd* argument –

<i>cmd</i> value	Meaning
<b>SIG_SETMASK</b>	Overrides the calling process signal mask with the value specified in the <i>new_mask</i> argument
<b>SIG_BLOCK</b>	Adds the signals specified in the <i>new_mask</i> argument to the calling process signal mask
<b>SIG_UNBLOCK</b>	Removes the signals specified in the <i>new_mask</i> argument from the calling process signal mask

- If the actual argument to *new\_mask* argument is a NULL pointer, *cmd* argument will be ignored and current process signal mask will not be altered.
- *old\_mask* argument – address of a *sigset\_t* variable that will be assigned the original signal mask of calling process prior to *sigprocmask* call.
- If actual argument to *old\_mask* is NULL pointer, no previous signal mask will be returned.

- *sigset\_t* is a data type defined in *<sigset.h>* - contains a collection of bit flags, with each bit flag representing one signal defined in a given system.
- *sigsetops* functions - set of API defined by BSD UNIX and POSIX.1 to set, reset and query the presence of signals in a *sigset\_t* typed variable.

```
#include<signal.h>

int sigemptyset(sigset_t *sigmask);
int sigaddset(sigset_t *sigmask, const int signal_num);
int sigdelset(sigset_t *sigmask, const int signal_num);
int sigfillset(sigset_t *sigmask);
int sigismember(const sigset_t *sigmask, const int signal_num);
```

- *sigemptyset* - clears all signal flags in *sigmask* argument.
- *sigaddset* - sets the flag corresponding to *signal\_num* signal in the *sigmask* argument.
- *sigdelset* - clears the flag corresponding to *signal\_num* signal in the *sigmask* argument.
- *sigfillset* - sets all signal flags in *sigmask* argument.
- *sigismember* - checks if the flag corresponding to *signal\_num* signal in the *sigmask* argument is set or not set.
- Return value of *sigemptyset*, *sigaddset*, *sigdelset* and *sigfillset* calls is 0 if the calls succeed or -1 if they fail.
- Possible causes of failure may be that the *sigmask* and/or *signal\_num* arguments are invalid.
- The *sigismember* API returns 1 if the flag corresponding to the *signal\_num* signal in the *sigmask* argument is set, 0 if it is not set, and -1 if the call fails.
- The following example checks whether the SIGINT signal is present in a process signal mask and adds it to the mask if it is not there.

Then it clears the SIGSEGV signal from the process signal mask.

```
#include<stdio.h>
#include<signal.h>

int main(){
    sigset_t sigmask;
    sigemptyset(&sigmask); //initialize set
    if(sigprocmask(0,0,&sigmask) == -1){ //get current signal mask
        perror("sigprocmask");
        exit(1);
    }
    else
        sigaddset(&sigmask, SIGINT); //set SIGINT flag
    sigdelset(&sigmask, SIGSEGV); //clear SIGSEGV flag
    if(sigprocmask(SIG_SETMASK,&sigmask,0) == -1)
        perror("sigprocmask"); //set new signal mask
}
```

- When one or more signals are pending for a process and are unblocked via *sigprocmask* API, the signal handler methods for those signals that are in effect at the time of *sigprocmask* call will be applied before the API is returned to caller.
- If multiple instances of the same signal are pending for the process, then it is implementation-dependent whether one or all the instances will be delivered to the process.
- A process can query which signals are pending for it using *sigpending* API:

```
#include<signal.h>

int sigpending(sigset_t *sigmask);
```

- *sigmask* – address of a *sigset\_t* typed variable; assigned the set of signals pending for the calling process by the API.
- The API returns 0 if it succeeds and -1 if it fails.
- The *sigpending* API can be useful to find out whether one or more signals are pending for a process and to set up signal handling methods for these signals before the process calls the *sigprocmask* API to unblock them.

- The following example reports to the console whether the SIGTERM signal is pending for the process:

```
#include<iostream.h>
#include<stdio.h>
#include<signal.h>

int main(){
    sigset_t sigmask;
    sigemptyset(&sigmask);
    if(sigpending(&sigmask) == -1)
        perror("sigpending");
    else
        cout << "SIGTERM signal is: "
              << (sigismember(&sigmask, SIGTERM)?"set":"no set");
}
```

- The signal mask manipulation functions are –

```
#include<signal.h>

int sighold(int signal_num);
int sigrelse(int signal_num);
int sigignore(int signal_num);
int sigpause(int signal_num);
```

- sighold* API – adds the named signal *signal\_num* to the calling process signal mask.
- It is the same as using the *sigset* API with the SIG\_HOLD action:  
`sigset(<signal_num>, SIG_HOLD);`
- sigrelse* API – removes the named signal *signal\_num* for the calling process signal mask.
- sigignore* API – sets the handling method for the named signal *signal\_num* to SIG\_DFL
- sigpause* API – removes the named signal *signal\_num* from the calling process signal mask and suspends the process until it is interrupted by a signal.

## (V) **sigaction**

- sigaction* is the replacement for *signal* API in the latest UNIX and POSIX systems.
- sigaction* API is called by a process to set up a signal handling method for each signal it wants to deal with.
- It passes back the previous signal handling method for a given signal.
- It blocks the signal it is catching, allowing a process to specify additional signals to be blocked when the API is handling a signal.



- Prototype of *sigaction* API is given by:

```
#include<signal.h>

int sigaction(int signal_num, struct sigaction *action,
              struct sigaction *old_action);
```

- *struct sigaction* data type is defined in <signal.h> as:

```
struct sigaction{
    void (*sa_handler)(int);
    sigset_t sa_mask;
    int sa_flag;
};
```

- *sa\_handler* – corresponds to second argument of *signal* API.
- Can be set to SIG\_IGN, SIG\_DFL, or user-defined signal handler function.
- *sa\_mask* – specifies additional signals that a process wishes to block when it is handling the *signal\_num* signal.
- It does not block the signals currently specified in the signal mask of the process and *signal\_num* signal.
- *signal\_num* - designates which signal handling action is defined in the *action* argument.
- Previous signal handling method for *signal\_num* will be returned via *old\_action* argument if it is not a NULL pointer.
- If *action* argument is a NULL pointer, existing signal handling method of the calling process for *signal\_num* will be unchanged.
- The following program illustrates the use of *sigaction*:

```
#include<iostream.h>
#include<stdio.h>
#include<unistd.h>
#include<signal.h>

void callme(int sig_num){
    cout << "catch signal: " << sig_num;
}

int main(int argc, char *argv[]){
    sigset_t sigmask;
    struct sigaction action, old_action;
    sigemptyset(&sigmask);
    if(sigaddset(&sigmask, SIGTERM) == -1 ||
       sigprocmask(SIG_SETMASK,&sigmask,0) == -1)
        perror("set signal mask");
    sigemptyset(&action.sa_mask);
    sigaddset(&action.sa_mask, SIGSEGV);
```

```

    action.sa_handler = callme;
    action.sa_flag = 0;
    if(sigaction(SIGINT,&action,&old_action) == -1)
        perror("sigaction");
    pause(); //wait for signal interruption
    cout << argv[0] << " exists";
    return 0;
}

```

- The process signal mask is set with the SIGTERM signal.
- The process then defines a signal handler for the SIGINT signal and also specifies that the SIGSEGV signal is to be blocked when the process is handling the SIGINT signal.
- The process then suspends its execution via the *pause* API.
- If the SIGINT signal is generated to the process, the kernel first sets the process signal mask to block the SIGTERM, SIGINT and SIGSEGV signals.
- It then arranges the process to execute the *callme* signal handler function.
- When the *callme* function returns, the process signal mask is restored to contain only the SIGTERM signal, and the process will continue to catch the SIGILL signal.
- *sa\_flag* - used to specify special handling for certain signals.
- POSIX.1 defines two values for *sa\_flag*: 0 and SA\_NOCLDSTOP.
- If *sa\_flag* is 0, kernel will send the SIGCHLD signal to the calling process whenever its child process is either terminated or stopped.
- If *sa\_flag* is SA\_NOCLDSTOP, kernel will send the SIGCHLD signal to the calling process whenever its child process is terminated, but not when it is stopped.
- Additional flags for the *sa\_flag* field in UNIX System V.4 implementation – SA\_RESETHAND and SA\_RESTART
- SA\_RESETHAND – If *signal\_num* is caught, *sa\_handler* is set to SIG\_DFL before the signal handler function is called.  
*signal\_num* will not be added to the process signal mask when the signal handler function is executed.
- SA\_RESTART – If a signal is caught while a process is executing a system call, kernel will restart the system call after the signal handler returns.  
If this flag is not set in the *sa\_flag* field, the system call will be aborted with a return of -1 after the signal handler returns and *errno* will be set to EINTR.

**(VI) SIGCHLD signal and the waitpid Function**

- When a child process terminates or stops, kernel will generate a SIGCHLD signal to its parent process.
- Three different events may occur depending on how the parent sets up the handling of the SIGCHLD signal.
- *Event 1* – parent accepts the default action of the SIGCHLD signal:
  - SIGCHLD signal does not terminate the parent process.
  - Affects only the parent process if it arrives at the same time when the parent process is suspended by the *waitpid* system call.
  - Parent process will be awakened, *waitpid* will return the child's exit status and PID to the parent, and the kernel will clear up the process table slot for child.
  - With this setup, parent can call *waitpid* repeatedly to wait for each child it created.
- *Event 2* – parent ignores the SIGCHLD signal:
  - SIGCHLD signal will be discarded.
  - Parent will not be disturbed even if it is executing the *waitpid* system call.
  - If parent calls the *waitpid* API, the API will suspend the parent until all its child processes have terminated.
  - Child process table slots will be cleared up by the kernel and API will return a value of -1 to the parent process.
- *Event 3* – parent catches the SIGCHLD signal:
  - Signal handler function will be called in the parent process whenever a child process terminates.
  - If SIGCHLD signal arrives while the parent process is executing a *waitpid* system call, the *waitpid* API may be restarted to collect the child exit status and clear its process table slot after the signal handler function returns.
  - API may be aborted, and the child process table slot is not freed, depending on the parent setup of the signal action for the SIGCHLD signal.
- Interaction between SIGCHLD and *wait* API is the same as that between SIGCHLD and *waitpid* API.
- Earlier versions of UNIX used the SIGCLD signal instead of SIGCHLD.
- SIGCLD signal is now obsolete, but most of the latest UNIX systems have defined SIGCLD to be the same as SIGCHLD for backward compatibility.

**(VII) sigsetjmp and siglongjmp APIs**

- *setjmp* and *sigsetjmp* – marks one or more locations in a user program.
- *longjmp* and *siglongjmp* – called to return to any of the marked locations.
- These APIs give the provision of inter-function *goto* capability.

- These are defined in POSIX.1 and on most UNIX systems that support signal masks.
- The function prototypes are:

```
#include<setjmp.h>

int sigsetjmp(sigjmpbuf env, int save_sigmask);
int siglongjmp(sigjmpbuf env, int ret_val);
```

- *sigsetjmp* and *siglongjmp* are created to support signal mask processing.
- It is implementation-dependent on whether a signal process mask is saved and restored when it invokes the *setjmp* and *longjmp* APIs.
- *sigsetjmp* – behaves similarly to the *setjmp* API, but it has a second argument *save\_sigmask*.
  - Allows a user to specify whether a calling process signal mask should be saved to the provided *env* argument.
  - If the *save\_sigmask* argument is non-zero, the caller's signal mask is saved. Otherwise, signal mask is not saved.
- *siglongjmp* – behaves similarly to the *longjmp* API, but it also restores a calling process signal mask if the mask was saved in its *env* argument.
  - *ret\_val* argument – specifies the return value of the corresponding *sigsetjmp* API when it is called by *siglongjmp*.
  - *ret\_val* value should be a non-zero number. If it is zero, *siglongjmp* API will reset it to 1.
- *siglongjmp* API is usually called from user-defined signal handling functions.
- This is because the process signal mask is modified when a signal handler is called, and *siglongjmp* should be called to ensure that the process signal mask is restored properly when jumping out from a signal handling function.
- The following program illustrates the uses of *sigsetjmp* and *siglongjmp* APIs.
- The program sets its signal mask to contain SIGTERM, then sets up a signal trap for the SIGINT signal.
- The program then calls *sigsetjmp* to store its code location in the *env* global variable.
- The *sigsetjmp* call returns a 0 value when it is called directly in the user program and not via *siglongjmp*.
- The program suspends its execution via the *pause* API.
- When a user interrupts the process from the keyboard, the *callme* function is called.
- The *callme* function calls the *siglongjmp* API to transfer program flow back to the *sigsetjmp* function in the *main* function, which returns a value of 2.

```

#include<iostream.h>
#include<stdio.h>
#include<unistd.h>
#include<signal.h>
#include<setjmp.h>

sigjmp_buf env;

void callme(int sig_num){
    cout << "catch signal: " << sig_num;
    siglongjmp(env, 2);
}

int main(){
    sigset_t sigmask;
    struct sigaction action, old_action;
    sigemptyset(&sigmask);
    if(sigaddset(&sigmask, SIGTERM) == -1 ||
        sigprocmask(SIG_SETMASK, &sigmask, 0) == -1)
        perror("set signal mask");
    sigemptyset(&action.sa_mask);
    sigaddset(&action.sa_mask, SIGSEGV);

```

### (viii) kill

- A process can send a signal to a related process using *kill* API.
- Sender and recipient processes must be related such that either the sender process real/effective UID matches that of recipient process, or sender process has superuser privileges.
- Example: parent and child process sending signals to each other.
- *kill* API function prototype:

```

#include<signal.h>

int kill(pid_t pid, int signal_num);

```

- Returns: 0 if successful, -1 if it fails.
- The following program illustrates the implementation of the UNIX *kill* command using the *kill* API.

```

#include<iostream.h>
#include<stdio.h>
#include<unistd.h>
#include<string.h>
#include<signal.h>

int main(int argc, char *argv[]){
    int pid, sig = SIGTERM;
    if(argc == 3){
        if(sscanf(argv[1], "%d", &sig) != 1){
            //get signal number
            cerr << "Invalid number: " << argv[1];
            return -1;
        }
        argv++, argc--;
    }
    while(--argc > 0){
        if(sscanf(*++argv, "%d", &pid) == 1){
            //get process ID
            if(kill(pid, sig) == -1)
                perror("kill");
        }
        else
            cerr << "Invalid pid" << argv[0];
    }
    return 0;
}

```

- *signal\_num* – integer value of a signal that must be sent to one or more processes designated by *pid*.
- Possible values of *pid*: positive value, 0, -1, negative value.
- The following table gives the details about these values:

<i>pid</i> value	Effects on <i>kill</i> API
<b>Positive value</b>	<i>pid</i> is a process ID. Sends <i>signal_num</i> to that process.
<b>0</b>	Sends <i>signal_num</i> to all processes whose process GID is same as the calling process
<b>-1</b>	Sends <i>signal_num</i> to all process whose real UID is same as the effective UID of calling process. If calling process effective UID is the superuser UID, <i>signal_num</i> will be sent to all processes in the system (except process 0 and process 1).
<b>Negative value</b>	Sends <i>signal_num</i> to all processes whose process GID matches absolute value of <i>pid</i>

- UNIX *kill* command invocation syntax is given as follows:  
*kill [-<signal\_num>] <pid>*
- Here, <signal\_num> - integer number or symbolic name of a signal, as defined in <signal.h> header.
- <pid> - integer number of a PID.  
There can be one or more PID specified – *kill* will send signal <signal\_num> to each process that corresponds to a <pid>.
- Example: *kill -9 1234*
- Any signal specification at the command line must be a signal's integer value.
- It does not support signal symbolic names.
- If no signal number is specified, the program will use the default signal SIGTERM, which is the same for the UNIX *kill* command.
- The program calls the *kill* API to send a signal to each process whose PID is specified at the command line.
- If a PID is invalid or if the *kill* API fails, the program will flag an error message.

## (IX) alarm

- *alarm* can be called by a process to request the kernel to send the SIGALRM signal after a certain number of real clock seconds.
- *alarm* API function prototype:

```
#include<signal.h>

unsigned int alarm(unsigned int time_interval);
```

- Returns: number of CPU seconds left in the process timer, as set by previous *alarm* system call.
- *time\_interval* – number of CPU seconds elapsed time, after which the kernel will send the SIGALRM signal to the calling process.  
If *time\_interval* = 0, it turns off the alarm clock.
- Effect of previous *alarm* is cancelled and process timer is reset with new *alarm* call.
- Process alarm clock is not passed on to its *forked* child, but an *execed* process retains the same alarm clock value as was prior to the *exec* API call.
- *alarm* API is used to implement *sleep* API – suspends a calling process for the specified number of CPU seconds.
- Process will be awakened by either the elapsed time exceeding the *timer* value or when the process is interrupted by a signal.

- The implementation is given below:

```
#include<signal.h>
#include<stdio.h>
#include<unistd.h>

void wakeup(){};

unsigned int sleep(unsigned int timer){
    struct sigaction action;
    action.sa_handler = wakeup;
    action.sa_flags = 0;
    sigemptyset(&action.sa_mask);
    if(sigaction(SIGALRM,&action,0) == -1){
        perror("sigaction");
        return -1;
    }
    (void)alarm(timer);
    (void)pause();
    return 0;
}
```

- The *sleep* function sets up a signal handler for SIGALRM, calls the *alarm* API to request the kernel to send the SIGALRM signal after the *timer* interval and suspends its execution via the *pause* system call.
- The *wakeup* signal handler function is called when the SIGALRM signal is sent to the process.
- When it returns, the *pause* system call will be aborted, and the calling process will return from the *sleep* function.
- BSD UNIX defines the *ualarm* function.
- It is same as that of *alarm* API, but argument and return values are in microseconds.
- This is particularly useful for time-critical applications where the resolution of time must be in microsecond levels.
- This can be used to implement BSD-specific *usleep* function, which is similar to *sleep*, but its argument is in microseconds.

## (X) Interval Timers

- Interval timers are implemented using the *alarm* API.
- These can be used to schedule a process to do some tasks at a fixed time interval, to time the execution of some operations, or to limit the time allowed for the execution of some tasks.



- The following program illustrates how to set up a real-time clock interval timer using the *alarm* API.

```
#include<stdio.h>
#include<unistd.h>
#include<signal.h>
#define INTERVAL 5

void callme(int sig_num){
    alarm(INTERVAL);
    //do scheduled tasks
}

int main(){
    struct sigaction action;
    sigemptyset(&action.sa_mask);
    action.sa_handler = (void (*)())callme;
    action.sa_flag = SA_RESTART;
    if(sigaction(SIGALRM,&action,0) == -1){
        perror("sigaction");
        return 1;
    }
    if(alarm(INTERVAL) == -1)
        perror("alarm");
    else
        while(1){
            //do normal operation
        }
    return 0;
}
```

- The *sigaction* API is called to set up *callme* as the signal handling function for the SIGALRM signal.
- The program then invokes the *alarm* API to send itself the SIGALRM after 5 real clock seconds.
- The program then goes off to perform its normal operation in an infinite loop.
- When the timer expires, the *callme* function is invoked, which restarts the alarm clock for another 5 seconds and then does the scheduled tasks.
- When the *callme* function returns, the program returns continues its normal operation until another timer expiration.
- *setitimer* API – additional capabilities defined in BSD UNIX. Also available in UNIX System V.3 and V.4, but not specified by POSIX.
- POSIX.1b has new set of APIs for interval timer manipulation.

- Additional features of *setitimer* API are:
  - *alarm* resolution time is in seconds.
  - setitimer* resolution time is in microseconds.
  - *alarm* can be used to set up 1 real-time clock timer per process.
  - setitimer* can set up 3 - real-time clock timer, timer based on user time spent by a process, timer based on total user time and system times spent by a process.
- setitimer* and *getitimer* function prototypes are given as follows:

```
#include<sys/time.h>

int setitimer(int which, const struct itimerval *val,
              struct itimerval *old);
int getitimer(int which, struct itimerval *old);
```

- Return values for both: 0 if successful, -1 if fail.
- Timers set by *setitimer* in a parent process will not be inherited by child processes, but are retained when a process *execs* a new program.
- which* argument – specifies which timer to process.
- The following table specifies the various values for *which* argument –

<i>which</i> value	Timer type
ITIMER_REAL	Timer based on real-time clock. Generates a SIGALRM signal when it expires.
ITIMER_VIRTUAL	Timer based on user time spent by a process. Generates SIGVTALRM signal when it expires.
ITIMER_PROF	Timer based on total user and system times spent by a process. Generates a SIGPROF signal when it expires.

- ITIMER\_VIRTUAL and ITIMER\_PROF timers – useful in timing the total execution time of selected user functions, as the timer runs only while the user process is running.
- struct itimerval* defined in *<sys/time.h>* header as –

```
struct itimerval{
    struct timeval it_interval; //timer interval
    struct timeval it_value; //current value
};
```

- For *setitimer* API – *it\_value* is the time to set the named timer, *it\_interval* is the time to reload the timer when it expires.
- For *getitimer* API – *it\_value* is the named timer's remaining time to expiration, *it\_interval* is the time to reload the timer when it expires.

- The following timer illustrates how to set up a real-time clock interval timer using the *setitimer* API.
- The timer is specified to be reloaded automatically.

```
#include<stdio.h>
#include<unistd.h>
#include<sys/time.h>
#include<signal.h>
#define INTERVAL 2

void callme(int sig_num){
    //do scheduled tasks
}

int main(){
    struct itimerval val;
    struct sigaction action;
    sigemptyset(&action.sa_mask);
    action.sa_handler = (void (*)())callme;
    action.sa_flag = SA_RESTART;
    if(sigaction(SIGALRM,&action,0) == -1){
        perror("sigaction");
        return 1;
    }

    val.it_interval.tv_sec = INTERVAL;
    val.it_interval.tv_usec = 0;
    val.it_value.tv_sec = INTERVAL;
    val.it_value.tv_usec = 0;
    if(setitimer(IIMER_REAL,&val,0) == -1)
        perror("alarm");
    else
        while(1){
            //do normal operation
        }
    return 0;
}
```

- The real time clock timer set by the *setitimer* API is different from that set by the *alarm* API.
- Thus, a process may set up two real-time clock timers using the two APIs.
- Furthermore, since the *alarm* and *setitimer* APIs require that users set up signal handling to catch timer expiration, they should not be used in conjunction with the *sleep* API.
- This is because the *sleep* API may modify the signal handling function for the SIGALRM signal.

## (XI) POSIX.1b Timers

- POSIX.1b timers are a set of APIs for interval timer manipulation defined in POSIX.1b implementation.
- These are more flexible and powerful than UNIX timers –
  - Users may define multiple independent timers per system clock.
  - Timer resolution is in nanoseconds.
  - Users may specify the signal to be raised when a timer expires, on a per-timer basis.
  - Timer interval may be specified as either an absolute or a relative time.
- The limit on maximum number of POSIX timers per process is given by `TIMER_MAX` constant defined in `<limits.h>` header.
- POSIX timers created by a process are not inherited by its child process, but are retained across the `exec` system call.
- Unlike UNIX timers, POSIX timer can be used safely with `sleep` API if it does not use the `SIGALRM` signal when it expires.
- The POSIX.1b APIs for timer manipulation are:

```
#include<signal.h>
#include<time.h>

int timer_create(clockid_t clock, struct sigevent *spec,
                 timer_t *timer_hdrp);
int timer_settime(timer_t timer_hdr, int flag, struct itimerspec *val,
                 struct itimerspec *old);
int timer_gettime(timer_t timer_hdr, struct itimerspec *old);
int timer_getoverrun(timer_t timer_hdr);
int timer_delete(timer_t timer_hdr);
```

- `timer_create` API – used to dynamically create a timer and return its handler.
- `clock` argument – specifies which system clock the new timer should be based on.
- `clock` value may be `CLOCK_REALTIME` for creating a real-time clock timer.
- `spec` argument – defines what action to take when the timer expires.
- The `struct sigevent` data type is defined as:

```
struct sigevent{
    int sigev_notify;
    int sigev_signo;
    union sigval sigev_value;
};
```

- The data structure of the `sigev_value` field is:

```
union sigval{
    int sival_int;
    void *sival_ptr;
};
```

- `timer_settime` API – used to start or stop a running timer.

- *flag* argument in *timer\_settime* API may be 0 or *TIMER\_RELTIME* if the timer start time is relative to the current time.
- If the *flag* argument is *TIMER\_ABSTIME*, the timer start time is an absolute time.
- The ANSI C *mktime* function may be used to generate the absolute time for setting a timer.
- If the *val.it\_value* is 0, it stops the timer from running.
- If the *val.it\_interval* is 0, the timer will not restart after it expires.
- The *old* argument of *timer\_settime* API is used to obtain the previous timer values.
- It may be set to NULL, and no timer values are returned.
- *timer\_gettime* API – used to query the current values of a timer.
- *old* argument of *timer\_gettime* API returns the current values of the named timer.
- The *struct itimerspec* data type is defined as:

```
struct itimerspec{
    struct timespec it_interval;
    struct timespec it_value;
};
```

- The *itimerspec::it\_value* specifies the time remaining in the timer, and the *itimerspec::it\_interval* specifies the new time to reload the timer after it expires.
- The *struct timespec* data structure is defined as:

```
struct timespec{
    time_t tv_sec;
    long tv_nsec;
};
```

- All times are specified in seconds via the *timespec::tv\_sec* field, and in nanoseconds via the *timespec::tv\_nsec* field.
- *timer\_getoverrun* API – returns the number of signals generated by a timer but was lost / overrun.
- Timer signals are not queued by the kernel if they are raised but are not being handled by their target processes.
- Instead, the kernel records the number of these overrun signals per timer.
- The *timer\_getoverrun* API can be used to determine the amount of time elapsed between the timer started or handled to the present time, based on the overrun count of a named timer.
- The overrun count in a timer is reset whenever a process handles the timer signal.
- *timer\_delete* API – used to destroy a timer created by *timer\_create* API.
- The following program illustrates how to set up an absolute-time timer that would have gone off at 23:59:59, on 31 December 1999.

```
#include<iostream.h>
#include<stdio.h>
#include<unistd.h>
#include<signal.h>
#include<time.h>
#define TIMER_TAG 12

void callme(int signo, siginfo_t *evp, void *ucontext){
    time_t tim = time(0);
    cerr << "callme: " << evp->si_value.sival_int
        << ", signo: " << signo << ", " << ctime(&tim);
}

int main(){
    struct sigaction sigv;
    struct sigevent sigx;
    struct itimerspec val;
    struct tm do_time;
    timer_t t_id;

    sigemptyset(&sigv.sa_mask);
    sigv.sa_flag = SA_SIGINFO;

    sigv.sa_sigaction = callme;
    if(sigaction(SIGUSR1, &sigv, 0) == -1){
        perror("sigaction");
        return 1;
    }
    sigx.sigev_notify = SIGEV_SIGNAL;
    sigx.sigev_signo = SIGUSR1;
    sigx.sigev_value.sival_int = TIMER_TAG;
    if(timer_create(CLOCK_REALTIME, &sigx, &t_id) == -1){
        perror("timer create");
        return 1;
    }
    //set timer to go off at December 31, 1999, 23:59:59
    do_time.tm_hour = 23;
    do_time.tm_min = 59;
    do_time.tm_sec = 59;
    do_time.tm_mon = 12;
    do_time.tm_year = 99;
    do_time.tm_mday = 31;

    val.it_value.tv_sec = mktime(&do_time);
    val.it_value.tv_nsec = 0;
```

```

    val.it_interval.tv_sec = 15;
    val.it_interval.tv_nsec = 0;
    cerr<<"Timer will go off at:"<<ctime(&val.it_value.tv_sec);

    if(timer_settime(t_id,TIMER_ABSTIME,&val,0) == -1){
        perror("timer_settime");
        return 2;
    }

    //do something, then wait for the timer to expire twice
    for(int i=0; i<2; i++)
        pause();
    if(timer_delete(t_id) == -1){
        perror("timer_delete");
        return 3;
    }
    return 0;
}

```

- The above program first sets up the *callme* function as the signal handler for the SIGUSR1 signal.
- It then creates a timer based on the system real-time clock.
- The program specifies that the timer should raise the SIGUSR1 signal whenever it expires, and the timer-specific data that should be sent along with the signal is *TIMER\_TAG*.
- The timer handler returned by the *timer\_create* API is stored in the *t\_id* variable.
- The next step is to set the timer to go off on 31 December 1999, at 23:59:59 hours and the timer should re-run for every 30 seconds thereafter.
- The absolute expiration date/time is specified in the *do\_time* variable (of type *struct tm*) and is being converted to a *time\_t*-type value via the *mktime* function.
- After all these are done, the *timer\_settime* function is called to start the timer running.
- The program then waits for the timer to expire at the said date/time and expires again 30 seconds later.
- Finally before the program terminates, it calls the *timer\_delete* to free all system resources allocated for the timer.
- The above program can be modified to use a relative-time timer instead, as follows. The *main* function is modified to set the timer to go off 60 minutes from now and repeat every 120 seconds after.

```
int main(){
    //set up sigaction for SIGUSR1
    ...
    //create a timer using timer_create
    ...
    struct itimerspec val;
    val.it_value.tv_sec = 60; //expire 60s from now
    val.it_value.tv_nsec = 0;
    val.it_interval.tv_sec = 120; //repeat every 120s
    val.it_interval.tv_nsec = 0;
    if(timer_settime(t_id,0,&val,0) == -1){
        perror("timer_settime");
        return 2;
    }
    //wait for timer to expire
    ...
}
```

- The differences between the *main* functions of this program and the previous program are:
  - the *do\_time* variable and *mktime* API are not being used.
  - the *val.it\_value* is set directly with the relative time (from the present) when the timer will first expire.
  - the second argument to *timer\_settime* call is set to 0 instead of *TIMER\_ABSTIME*.

## (XII) Daemons

- Daemons are processes that live for a long time.
- These are often started when system is bootstrapped and terminate only when the system is shut down.
- There is no controlling terminal for daemons; they always run in the background.
- Perform day-to-day activities in the UNIX system.
- Process names of a daemon usually end with "d".
- Examples: *syslogd* – a daemon that implements system logging facility, *sshd* – a daemon that serves incoming SSH connections.
- Daemons respond to network requests, hardware activity, or other programs by performing some task.
- *cron* daemon performs defined tasks at scheduled times.
- Alternative terms for daemons – service (Windows and later versions of Linux), started task (IBM z/OS), ghost job (XDS UTS).
- Network services – daemons that connect to a computer network.



**(XIII) Daemon Characteristics**

- *ps* command – prints the status of various processes in the system.
- Options:
  - a: displays the status of process owned by others
  - x: displays the processes that do not have a controlling terminal
  - j: displays the job-related information – session ID, process group ID, controlling terminal and terminal process group ID.
- Typical output will contain – parent process ID, process ID, process group ID, session ID, terminal name, terminal process group ID, user ID and command string.
- Kernel processes – any process whose parent process ID is 0.
- These processes are special and generally exist for the entire lifetime of the system.
- These processes run with superuser privileges and have no controlling terminal and no command line.
- Process 1 – *init* process.
- It is the system daemon responsible for starting system services specific to various run levels.
- These services are usually implemented with the help of their own daemons.
- *keventd* daemon – provides process context for running scheduled functions in the kernel.
- *kapmd* daemon – provides support for the advanced power management features available with various computer systems.
- *kswapd* daemon – supports the virtual memory subsystem by writing dirty pages to disk slowly over time, so the pages can be reclaimed. (a.k.a pageout daemon)
- *bdflush* daemon – flushes dirty buffers from the buffer cache back to disk when available memory reaches a low-water mark.
- *kupdated* daemon – flushes dirty pages back to disk at regular intervals to decrease data loss in the event of a system failure.
- *portmap* daemon – portmapper daemon which provides the service of mapping RPC (Remote Procedure Call) program numbers to network port numbers.
- *syslogd* daemon – available to any program to log system messages for an operator. The messages may be printed on a console device and written to a file.
- *inetd* daemon – listens on the system's network interfaces for incoming requests for various network servers.
- *nfsd*, *lockd*, *rpciod* daemons – provide support for the Network File System (NFS).

- *cron* daemon – executes commands at specified dates and times. Numerous system administration tasks are handled by having programs executed regularly by *cron*.
- *cupsd* daemon – handles print requests on the system (print spooler).
- Most of the daemons run with superuser privilege (a user ID of 0).
- None of the daemons has a controlling terminal: the terminal name is set to a question mark, and the terminal foreground process group is 1.
- All the user-level daemons are process group leaders and session leaders and are the only processes in their process group and session.
- The parent of most of these daemons is the *init* process.

#### (XIV) Coding Rules

- Certain rules are written to code a daemon so as to prevent unwanted interactions from happening in the system.
- *Rule 1* – Call *umask* to set the file mode creation mask to 0.
- Inherited file mode creation mask is set to deny certain permissions.
- If the daemon process is going to create files, the user may want to set specific permissions.
- Example: if a process specifically creates files with group-read and group-write permissions enabled, the file mode creation mask should be set accordingly.
- *Rule 2* – Call *fork* and make the parent *exit*.
- If the daemon was started as a simple shell command, having the parent terminate makes the shell think that the command is done.
- The child inherits the process group ID of the parent but gets a new process ID, which implies that the child is not a process group leader.
- This rule is a prerequisite for the call to *setsid* that is done next.
- *Rule 3* – Call *setsid* to create a new session.
- The process becomes a session leader of a new session.
- The process becomes the process group leader of a new process group.
- The process has no controlling terminal.
- *Rule 4* – Change the current working directory to the root directory.
- Current working directory inherited from parent could be on a mounted file system.
- Daemons normally exist until the system is rebooted.  
If it stays on a mounted file system, that file system cannot be unmounted.

- Alternate – some daemons might change the current working directory to some specific location, where they will do all their work.
- Example: line printer spooling daemons often change to their spool directory.
- *Rule 5* – Unneeded file descriptors should be closed.
- This rule prevents the daemon from holding open any descriptors that it may have inherited from its parent, which could be a shell or some other process.
- *open\_max* function or *getrlimit* function determine the highest descriptor and close all descriptors up to that value.
- *Rule 6* – Open file descriptors 0, 1, and 2 to */dev/null* so that any library routines that try to read from standard input or write to standard output or standard error will have no effect.
- As daemon is not associated with a terminal device, there is nowhere for output to be displayed; nor is there anywhere to receive input from an interactive user.
- Even if the daemon was started from an interactive session, the daemon runs in the background, and the login session can terminate without affecting the daemon.
- If other users log in on the same terminal device, output from the daemon should not show up on the terminal, and input from users should not be read by daemon.
- Based on the coding rules stated above, the following function is used to “daemonize” a process:

```
#include "apue.h"
#include<syslog.h>
#include<fcntl.h>
#include<sys/resource.h>

void daemonize(const char *cmd){
    int i, fd0, fd1, fd2;
    pid_t pid;
    struct rlimit rl;
    struct sigaction sa;

    //clear file creation mask
    umask(0);
    //get maximum number of file descriptors
    if(getrlimit(RLIMIT_NOFILE,&rl) < 0)
        err_quit("%s: cannot get file limit",cmd);
    //become session leader to lose controlling TTY
    if((pid=fork()) < 0)
        err_quit("%s: cannot fork",cmd);
    else if(pid != 0) //parent
        exit(0);
    setsid();
```

```

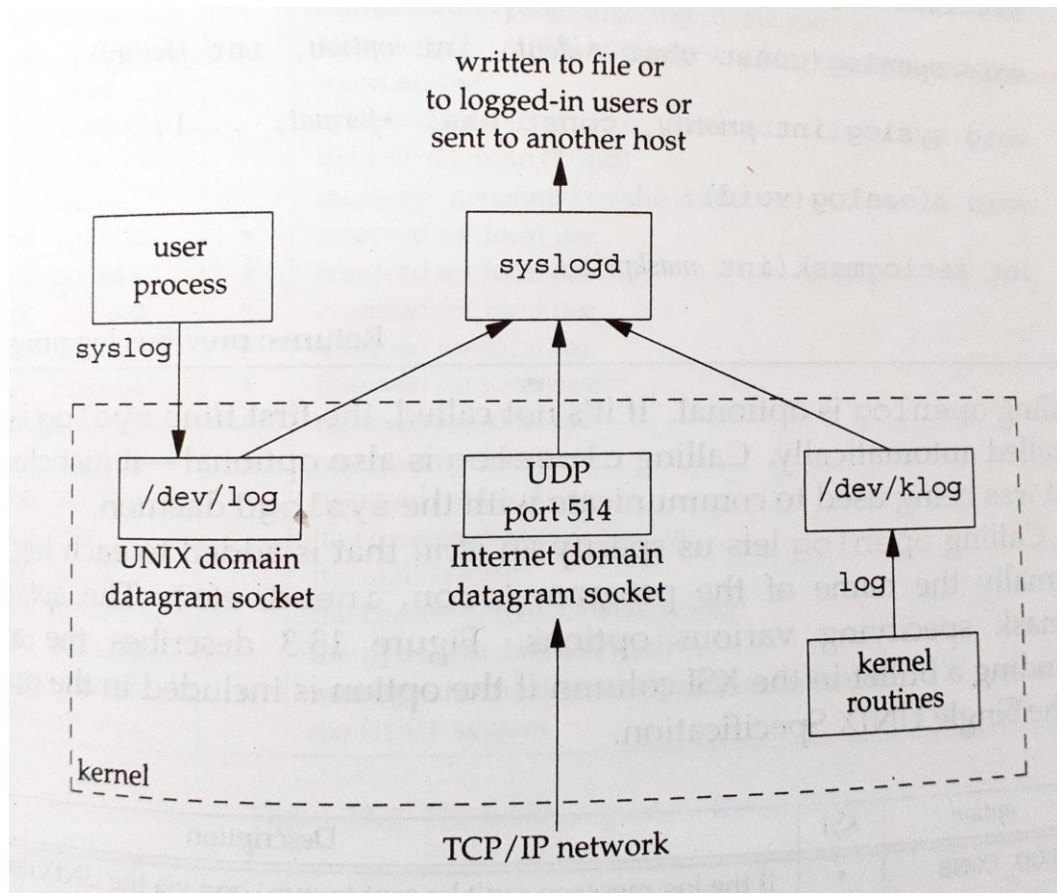
//ensure future opens will not allocate controlling TTYS
sa.sa_handler = SIG_IGN;
sigemptyset(&sa.sa_mask);
sa.sa_flag = 0;
if(sigaction(SIGHUP, &sa, NULL) < 0)
    err_quit("%s: cannot ignore SIGHUP");
if((pid=fork()) < 0)
    err_quit("%s: cannot fork",cmd);
else if(pid != 0) //parent
    exit(0);
//change current working directory to the root
if(chdir("/") < 0)
    err_quit("%s: cannot change directory to /");
//close all open file descriptors
if(rlim_max == RLIM_INFINITY)
    rlim_max = 1024;
for(i=0; i<rlim_max; i++)
    close(i);
//attach file descriptors 0,1 and 2 to /dev/null
fd0 = open("/dev/null", O_RDWR);
fd1 = dup(0);
fd2 = dup(0);
//initialize log file
openlog(cmd,LOG_CONS,LOG_DAEMON);
if(fd0 != 0 || fd1 != 1 || fd2 != 2){
    syslog(LOG_ERR,"unexpected file descriptors %d %d %d",
        fd0, fd1, fd2);
    exit(1);
}
}

```

#### (XV) Error Logging

- The main problem faced by a daemon is the handling of error messages.
- Daemon cannot write to standard error, as it should not have a controlling terminal.
- Daemons should not write to console device, as many workstations have windowing system running on console device.
- Daemons should not write error messages into a separate file, as it is impossible for admin to keep up with which daemon writes to which log file, and to check these logs on a regular basis.
- Thus, a central daemon error-logging facility is required.
- *syslog* facility developed at Berkeley, for the BSD 4.2 is used for error logging.
- Most systems derived from BSD support *syslog*.
- It is included as an XSI extension in Single UNIX Specification.
- Most daemons use *syslog* to log their error messages.

- The structure of *syslog* facility is given below –



- There are three ways to generate log messages:
  - Kernel routines can call the *log* function. These messages can be read by any user process that opens and reads the */dev/klog* device.
  - Most user processes (daemons) call the *syslog* function to generate log messages. This causes the message to be sent to the UNIX domain datagram socket */dev/log*.
  - A user process on this host, or on some other host that is connected to this host by a TCP/IP network, can send log messages to UDP port 514.
- syslog* never generates these UDP datagrams: they require explicit network programming by the process generating the log message.
- syslogd* daemon reads all three forms of log messages.
- On start-up, *syslogd* reads a configuration file, usually */etc/syslog.conf*, which determines where different classes of messages are to be sent.
- Example: urgent messages can be sent to the system administrator (if logged in) and printed on the console, whereas warnings may be logged to a file.

- Interface to the logging facility is through these functions:

```
#include<syslog.h>

void openlog(const char *ident, int option, int facility);
void syslog(int priority, const char *format, ...);
void closelog(void);
int setlogmask(int maskpri);
```

- Calling *openlog* is optional.  
If it is not called, when *syslog* is called for the first time, *openlog* is called automatically.
- Calling *closelog* is also optional.  
It just closes the descriptor that was being used to communicate with the *syslogd* daemon.
- ident* – added to each log message.  
This is normally the name of the program (*cron*, *inetd*, etc.)
- option* – bit mask specifying various options.
- The various options are listed in the table below:

<i>option</i>	Description
<b>LOG_CONS</b>	If the log message cannot be sent to <i>syslogd</i> via the UNIX domain datagram, the message is written to the console instead.
<b>LOG_NDELAY</b>	Open the UNIX domain datagram socket to the <i>syslogd</i> daemon immediately; do not wait until the first message is logged.
<b>LOG_ODELAY</b>	Delay the open of the connection to the <i>syslogd</i> daemon until the first message is logged.
<b>LOG_NOWAIT</b>	Do not wait for child processes that might have been created in the process of logging the message.
<b>LOG_PERROR</b>	Write the log message to standard error in addition to sending it to <i>syslogd</i> .
<b>LOG_PID</b>	Log the process ID with each message. This is intended for daemons that fork a child process to handle different requests.

- facility* – lets the configuration file specify that messages from different facilities are to be handled differently.
- If we do not call *openlog*, or if we call it with a *facility* of 0, we can still specify the facility as part of the *priority* argument to *syslog*.
- Single UNIX Specification defines only a subset of the facility codes typically available on a given platform.

- The following table gives details about the available facility codes:

<i>facility</i>	<b>Description</b>
<b>LOG_AUTH</b>	Authorization programs such as <i>login</i> , <i>su</i> , <i>getty</i>
<b>LOG_AUTHPRIV</b>	Same as LOG_AUTH, but logged to file with restricted permissions
<b>LOG_CRON</b>	<i>cron</i> and <i>at</i>
<b>LOG_DAEMON</b>	System daemons such as <i>inetd</i> , <i>routed</i>
<b>LOG_FTP</b>	FTP daemon <i>ftpd</i>
<b>LOG_KERN</b>	Messages generated by the kernel
<b>LOG_LOCAL0</b>	Reserved for local use
<b>LOG_LOCAL1</b>	Reserved for local use
<b>LOG_LOCAL2</b>	Reserved for local use
<b>LOG_LOCAL3</b>	Reserved for local use
<b>LOG_LOCAL4</b>	Reserved for local use
<b>LOG_LOCAL5</b>	Reserved for local use
<b>LOG_LOCAL6</b>	Reserved for local use
<b>LOG_LOCAL7</b>	Reserved for local use
<b>LOG_LPR</b>	Line printer system functions such as <i>lpd</i> , <i>lpc</i>
<b>LOG_MAIL</b>	Mail system
<b>LOG_NEWS</b>	Usenet network news system
<b>LOG_SYSLOG</b>	<i>syslogd</i> daemon itself
<b>LOG_USER</b>	Messages from other user processes
<b>LOG_UUCP</b>	UUCP system

- level* – ordered by priority, from highest to lowest.
- The following table lists the *level* values:

<i>level</i>	<b>Description</b>
<b>LOG_EMERG</b>	Emergency – system is unusable (highest priority)
<b>LOG_ALERT</b>	Condition that must be fixed immediately
<b>LOG_CRIT</b>	Critical condition (such as hard device error)
<b>LOG_ERR</b>	Error condition
<b>LOG_WARNING</b>	Warning condition
<b>LOG_NOTICE</b>	Normal, but significant condition
<b>LOG_INFO</b>	Informational message
<b>LOG_DEBUG</b>	Debug message (lowest priority)

- *priority* – combination of *facility* and *level*
- The *format* argument and any remaining arguments are passed to the *vsprintf* function for formatting.
- Any occurrence of the two characters *%m* in the *format* are first replaced with the error message string (*strerror*) corresponding to the value of *errno*.
- The *setlogmask* function can be used to set the log priority mask for the process.
- This function returns the previous mask.
- When the log priority mask is set, messages are not logged unless their priority is set in the log priority mask.
- Attempts to set the log priority mask to 0 will have no effect.
- *logger* program – provided by many systems to send log messages to the *syslog* facility.
- Optional arguments – specifies the *facility*, *level*, and *ident*.
- Single UNIX Specification does not define any options.
- *logger* is intended for a shell script running non-interactively that needs to generate log messages.
- Example: in a line printer spooler daemon, the following sequence might be seen:

```
openlog("lpd", LOG_PID, LOG_LPR);
syslog(LOG_ERR, "open error for %s: %m", filename);
```
- The first call sets the *ident* string to the program name, specifies that the process ID should always be printed, and sets the default facility to the line printer system.
- The call to *syslog* specifies an error condition and a message string.
- If *openlog* was not called, the second call could have been

```
syslog(LOG_ERR|LOG_LPR, "open error for %s: %m", filename);
```
- *vsyslog* – a variant of *syslog* that handles variable argument lists.

```
#include<syslog.h>
#include<stdarg.h>

void vsyslog(int priority, const char *format, va_list arg);
```
- However, this function is not included in the Single UNIX Specification.
- Most *syslogd* implementations will queue messages for a short time.
- If a duplicate message arrives during this time – the *syslog* daemon will not write it to the log.
- Instead, the daemon will print out a message similar to "last message repeated N times."



**(XVI) Client-Server Model**

- Common use for a daemon process is to be a server process.
- *syslogd* process – a server that has messages sent to it by user processes (clients) using a UNIX domain datagram socket.
- Server – a process that waits for a client to contact it, requesting some type of service.
- The service being provided by the *syslogd* server is the logging of an error message.
- Communication between the client and the server is one-way.  
The client sends its service request to the server; the server sends nothing back to the client.